


Desirable Properties

1. **Timeliness.** Results have to be correct not only in their value but also in time.
2. **Design for peak load.** Real-time systems must not collapse when they are subject to peak-load conditions.
3. **Predictability.** To guarantee a minimum level of performance, the system must be able to predict the consequences of any scheduling decision.
4. **Fault tolerance.** Single hardware and software failures should not cause the system to crash.
5. **Maintainability.** Possible system modifications are easy to perform.



Task Scheduling Concept

1. Many real-time systems have many tasks to be executed in parallel.
 2. An *efficient* scheduling algorithm is used to schedule all tasks that have some jobs ready to run.
- A single task queue for all jobs that are ready.
 - Tasks may *be assigned* to any processor for execution.
 - The **BIG** question is how we put jobs in front of the queue.

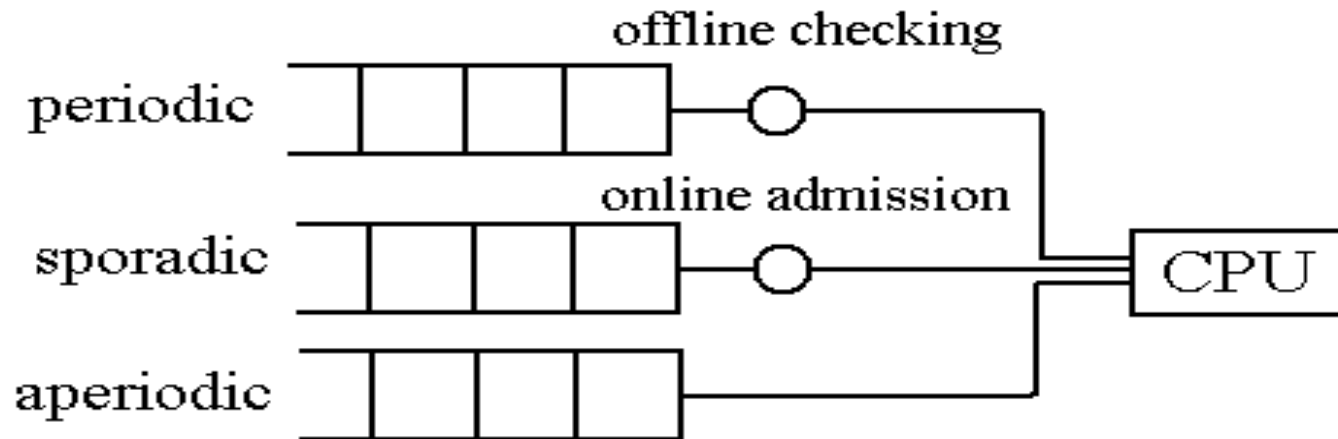


Real-time Task Models

- ▶ The term ‘**task**’ is a design-time concept. Each task represents a unit of concurrency. A task with real-time constraints is called a real-time task.
- ▶ A task can be executed multiple times. Each individual run to completion of a task is called a **job** of that task.
- ▶ **Periodic** tasks: a periodic task can be executed repeatedly on a continuing basis. The inter-arrival times between consecutive jobs are almost the same, which is called its period.
- ▶ **Sporadic** tasks: a sporadic task is executed in response to external events which occur at random instants of time. The inter-arrival times between consecutive jobs may vary widely, and can be arbitrarily small (bursty occurrence).
- ▶ **Aperiodic** tasks: an aperiodic task is also executed in response to external events which may occur at random instants of time. The inter-arrival times of aperiodic jobs may follow some probability distribution function. Aperiodic tasks have either soft or no deadlines.



Scheduling Classes



- ▶ We could use different schedulers in one real-time system for different classes of tasks

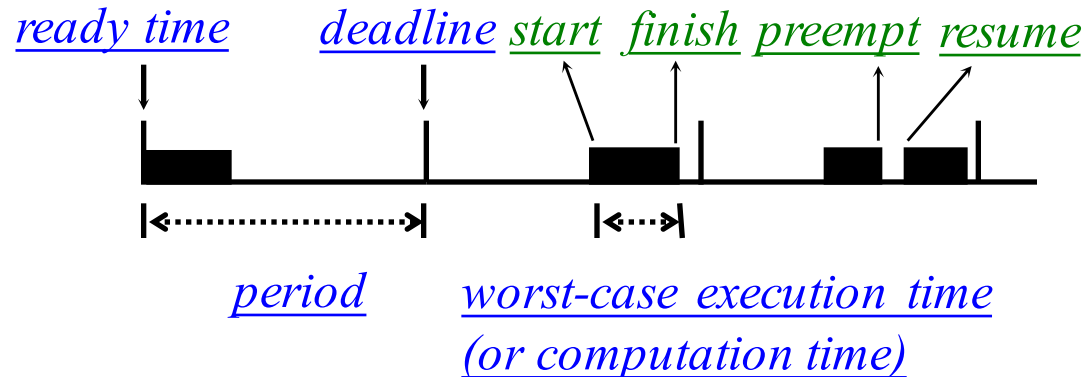
Model: Sampling Rate Decides Period

- ▶ **Sampling rate**
 - ▶ An application specification (e.g. 30 frames/sec, 6000 RPM)
- ▶ **Period**
 - ▶ How often a function is executed
- ▶ **Relative deadline**
 - ▶ The deadline with respect to when the function is ready for execution.



Periodic Job Scheduling

- ▶ Periodic Task Model
- ▶ parameters are known *a priori*



- Predictability vs. Schedulability
 - high predictability and flexibility
 - easy-to-check schedulability condition




A timing diagram showing the design of task schedule

Task A: period=20ms, execution time=6ms, deadline=20ms

Task B: period=30ms, execution time=10ms, deadline=30ms

Task C: period=40ms, execution time=3ms, deadline=40ms



Periodic Task vs. Job

- ▶ A periodic task is a stream of jobs (execution time, period)
- ▶ $T_1 = (3, 10) = \{J_{11} \ J_{12} \ J_{13} \ \dots\}$
- ▶ The time each job is eligible for execution is:
 - ▶ $(0, 10]$ $(10, 20]$ $(20, 30]$




Essential Real-Time Attributes

1. Deadline
2. Consequence (criticality)
 - ▶ Hard vs. Soft
3. Usefulness
4. Execution time



Deadline

- ▶ If a task's period is 30 ms, and the *relative* deadline is 10 ms, then the *absolute* deadlines are 10 ms, 40 ms, 70 ms...




- ▶ The relative deadline cannot be less than the computation, or execution time




Criticality

- ▶ A **hard** deadline must be completed, or a serious consequence will follow
 - ▶ Airplane instruments computing altitude
 - ▶ Air-bags in cars
- ▶ A **soft** deadline may be missed
 - ▶ How long are you willing to wait for your taxi reservation?



Value

- ▶ Most systems are a step function of value, with 1 for meeting the deadline and 0 for failure.




- ▶ Value of the success for a **hard** deadline is 0 or 1



Value (cont.)

- ▶ A **soft** deadline can have a range of values.
 - ▶ Completing before the deadline may not be advantageous.



Execution Time Issues

- ▶ **Execution time is not constant**
 - ▶ Load on the system
 - ▶ Interrupts and exception handling
 - ▶ Processor speed
 - ▶ OS/middleware overheads
 - ▶ Architecture: Cache, Pipeline, etc
- ▶ **Execution time has a distribution**
- ▶ **We assume there is a maximum execution time**



Worst-Case Task Execution Time

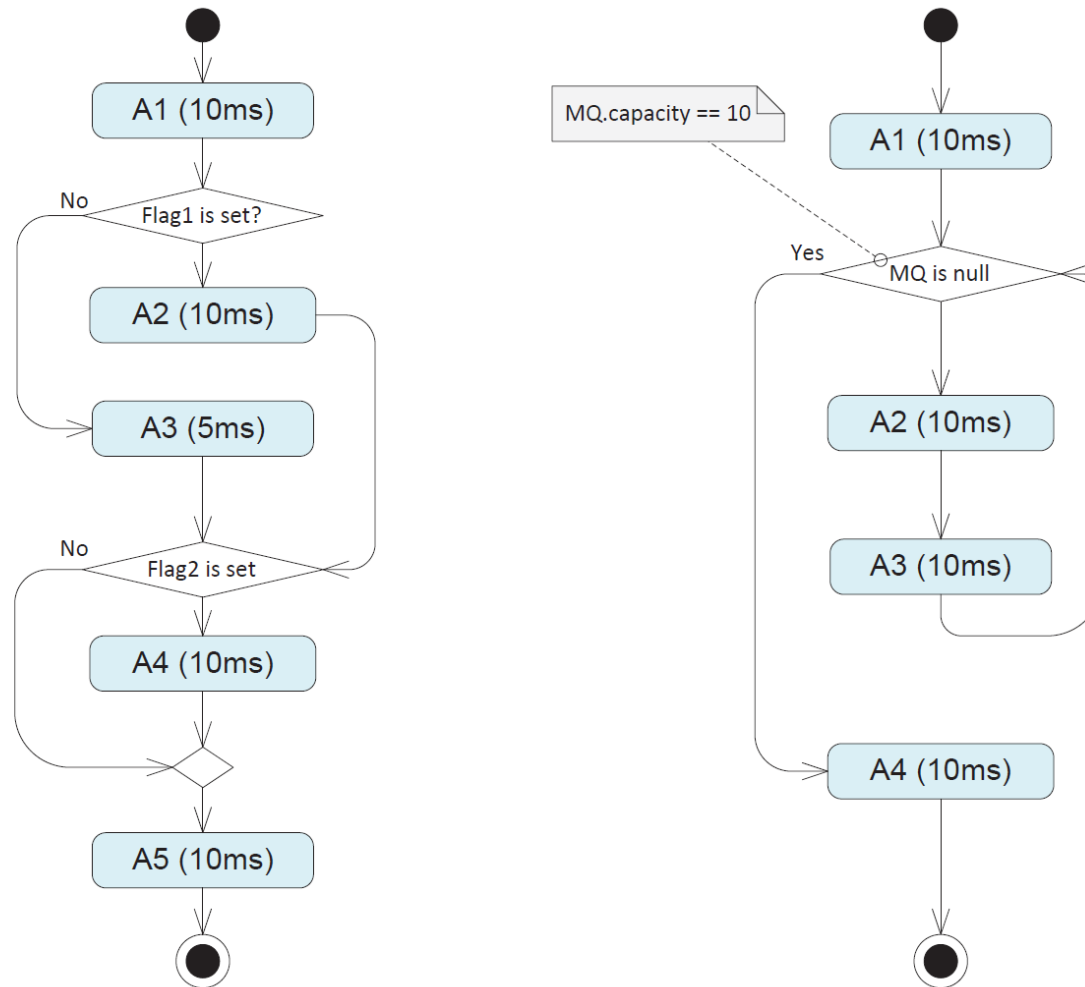


Fig. 12.1 The task shown on the left has a worst-case execution time of 40 ms, the task shown on the right has a worst-case execution time of 220 ms.

Study of Scheduling Algorithms

- ▶ Many process scheduling algorithms can be used by an OS kernel
 - ▶ FCFS, RR, EDF, Priority Driven...
- ▶ A kernel runs a variety of applications, and the performance of the scheduling algorithm depends on the application
 - ▶ Benchmark, Simulation
 - ▶ Analysis- Queuing theory gives averages, but can't guarantee a deadline
 - ▶ Validation- find the worst-case response time



Optimality

- ▶ Validation is a test that returns true or false
- ▶ “Optimality”
 - ▶ Create a set of tasks, and see which algorithm returns the most number of TRUE's
- ▶ If G_1 passes $S_1 = \{S_i\}$ and G_2 passes $S_2 = \{S_i\}$, then G_2 is better than G_1 if $S_2 \supset S_1$.




Algorithm Classification (I)

- ▶ **Preemptive.** With preemptive algorithms, the running task can be interrupted at any time to assign the processor to another active task, according to a predefined scheduling policy.
- ▶ **Non-preemptive.** With non-preemptive algorithms, a task, once started, is executed by the processor until completion. In this case, all scheduling decisions are taken as a task terminates its execution.




Preemptive Execution

- ▶ Preemptive execution: a job execution can be preempted by other jobs
 - ▶ The CPU is usually preemptive
 - ▶ A resource such as a printer usually is not.
- ▶ Phase: the first time a task becomes ready
 - ▶ Phase = 3:



Non-Preemptive Scheduling

- ▶ **Non-preemptive scheduling is risky**
 - ▶ Like checkout lines in a supermarket, you may wait for a long time if you don't know how long the person in front of you will take
 - ▶ $J_{12} = (\text{arrival}=2, \text{deadline}=4, \text{computation time} = 1)$
 - ▶ $J_{21} = (\text{arrival}=1.5, \text{deadline}=7, \text{computation time}=3)$



Algorithm Classification (II)

- ▶ **Static.** Static algorithms are those in which scheduling decisions are based on fixed parameters, assigned to tasks *before* their activation.
- ▶ **Dynamic.** Dynamic algorithms are those in which scheduling decisions are based on dynamic parameters that may change *during* system execution.



Algorithm Classification (III)

- ▶ **Online.** We say that a scheduling algorithm is used online if scheduling decisions are taken at *runtime* every time a new task enters the system or when a running task terminates.
- ▶ **Offline.** We say that a scheduling algorithm is used offline; if it is executed on the entire task set *before* actual task activation. The schedule generated in this way is stored in a table and later executed by a dispatcher.



Algorithm Classification (IV)

- ▶ **Optimal.** An algorithm is said to be optimal if it minimizes some cost function defined over the task set. When no cost function is defined and the only concern is to achieve a feasible schedule, then an algorithm is said to be optimal if it always finds a feasible schedule *whenever one exists*
- ▶ **Heuristic.** An algorithm is said to be heuristic if it searches for a feasible schedule using an objective function (*heuristic function*). Heuristic algorithm *do not guarantee* to find the optimal schedule, even if there exists one



Algorithm Classification (V): 1 vs. Multi-processor

- ▶ Migration: if a job is waiting for a resource and goes back to the queue, when it comes time to run the job again, it can go to any processor.
- ▶ Parallel execution: can parts of the job be executed simultaneously?
- ▶ Typically, allow migration but not parallel execution

