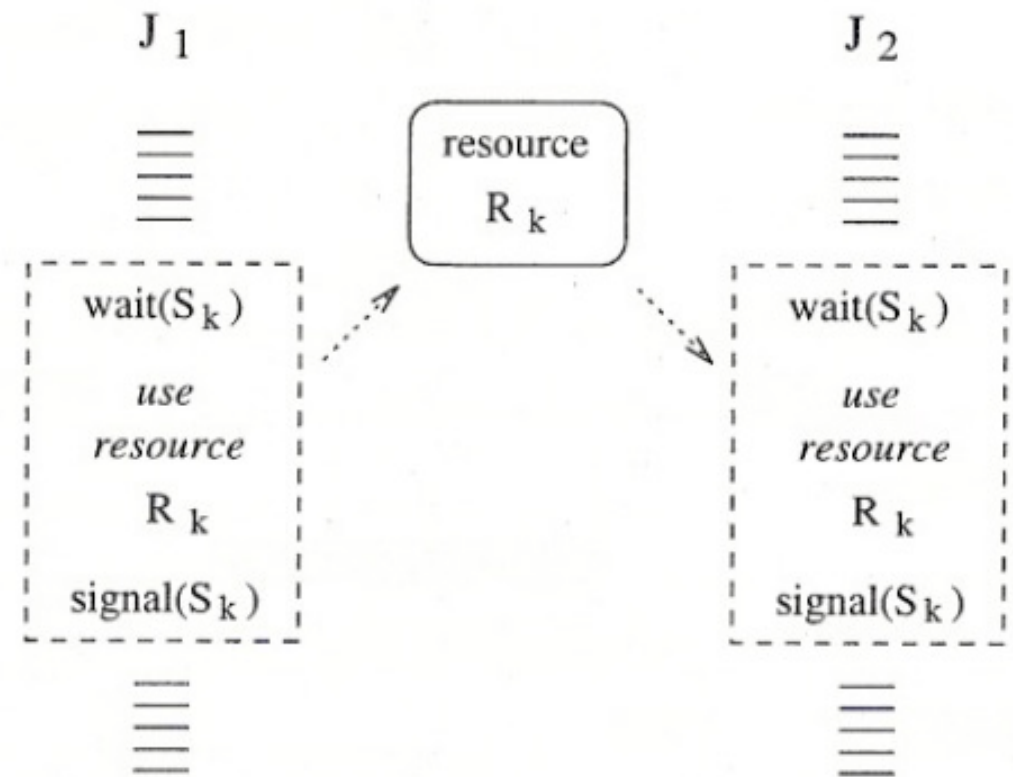


Introduction to Real-Time Computer Systems

Lecture 9

Sharing Resources

- ▶ Many real-time tasks must access resources (data, devices, files, etc.) that are shared among them.
- ▶ The resources are locked for integrity before usage.
- ▶ Low priority job may lock it first. And then high priority job will have to wait for the resource, causing extra delays.



Mars Rover Priority Inversion

- ▶ **Priority Inversion on Mars**

http://research.microsoft.com/en-us/um/people/mbj/mars_pathfinder/

http://research.microsoft.com/en-us/um/people/mbj/mars_pathfinder/Authoritative_Account.html

- ▶ **YouTube**

<http://www.youtube.com/watch?v=lyx7kARrGeM>


- ▶ **Detailed Technical Description**

http://retis.sssup.it/~marko/2009/mars_explorer.pdf




Priority Inversion, 1st case

- ▶ Assumption: Whoever locks a file for performing a write, no one else should be able to access the file.
 - ▶ Suppose J_1 and J_2 modify the same file.
 - ▶ J_1 has a higher priority than J_2 .
 - ▶ J_1 will be delayed by J_2 , between t_1 and t_2



Priority Inversion, 2nd case

- ▶ J_1 must wait for J_3 to release the lock before it can continue.
- ▶ Normally, J_1 is taking out of system contention. If the system then finds the next highest priority job and it's not J_3 ...



Priority Inversion, Problem Model

▶ Assumptions

- ▶ We know **what resources** will be used by each job
- ▶ Each resource has only one copy
- ▶ We don't know exactly when resources are locked
- ▶ There is a **maximum** length of time that resources are locked


▶ Problem

- ▶ Given a set of real-time jobs with all above parameters, can they be scheduled to meet deadline using RM or EDF?



Priority Inversion Solution I: NPP

- ▶ Non-Preemptive Protocol: When you lock a resource, you become the **highest** priority in the scheduler.
- ▶ Similar to the OS Kernel mode.



- ▶ Still delays J_1 , but not as much (no middle priority job can execute)
- ▶ Simple, but causes unnecessary delay



Highest Locker Protocol (HLP)

- ▶ A semaphore is given an “active priority” when it is created; the programmer should choose the highest priority of any thread that may **attempt** to lock the semaphore .
- ▶ As soon as a thread enters synchronized code, its (active) priority is **raised to the semaphore’s priority**.
 - ▶ If, through programming error, a thread has a higher base priority than the ceiling of the semaphore it is attempting to enter, then an exception is thrown.



Highest Locker Protocol (HLP)

- ▶ On leaving the semaphore, the thread has its active priority reset. In simple cases it will set to the thread's previous active priority, but under some circumstances (e.g. a dynamic change to the thread's base priority while it was in the semaphore) a different value is possible
- ▶ Benefit: A high priority job will not be affected by those critical sections shared only by low priority jobs.




Priority Inheritance Protocol (PIP)

- ▶ Whenever a job J_2 locks a resource,
 - ▶ Takes note of it but doesn't change priority.
 - ▶ When another task J_1 requests the locked resource, J_2 **inherits** (increases) to J_1 's priority if J_2 was lower.
 - ▶ J_1 goes to the ready queue.
 - ▶ When J_2 has unlocked the resource, it resumes the normal priority.




Example of PIP

■ normal execution
■ critical section



Two Kinds of Blocking



Two Kinds of Blocking

- ▶ **Direct blocking.** It occurs when a high-priority job tries to acquire a resource already held by a lower-priority job.
 - ▶ Direct blocking is **necessary** to ensure the consistency of the shared resources .
- ▶ **Push-through blocking.** It occurs when a medium-priority job is blocked by a lower-priority job that has inherited a higher priority from a job it directly blocks.
 - ▶ Push-through blocking is to **avoid** unbounded priority inversion.



Priority Inheritance Blocking Time

- ▶ Whenever a high priority job wants a resource locked by a lower priority job, the lower priority job gains the priority of the higher priority job as long as the higher priority job is waiting.
- ▶ The high priority job only needs to wait for one additional blocking time.
- ▶ So, we need to know the maximum time any lower priority job will hold a lock on the resource.



Blocking Properties

- ▶ **Lemma 7.1** A semaphore S_k can cause push-through blocking to job J_i , only if S_k is accessed both by a job with priority lower than P_i and by a job that has or can inherit a priority equal to or higher than P_i .
- ▶ **Lemma 7.3** If there are n lower-priority jobs that can block a job J_i , then J_i can be blocked for at most the duration of n critical sections (one for each of the n lower priority jobs), regardless of the number of semaphores used by J_i .
 - ▶ A job can block another if it is inside a critical section. Once out, it can't enter another since it no longer has a high priority.
- ▶ **Lemma 7.4** If there are m distinct semaphores that can block a job J_i , then J_i can be blocked for at most the duration of m critical sections, one for each of the m semaphores.
 - ▶ A nested critical section uses the outer most duration since it is the longest.
- ▶ **Theorem 7.1 (Sha-Rajkumar-Lehoczky)** Under the Priority Inheritance Protocol, a job J can be blocked for at most the duration of $\min(n, m)$ critical sections, where n is the number of lower-priority jobs that could block J and m is the number of distinct semaphores that can be used to block J .

