


Two Kinds of Blocking



Two Kinds of Blocking

- ▶ **Direct blocking.** It occurs when a high-priority job tries to acquire a resource already held by a lower-priority job.
 - ▶ Direct blocking is **necessary** to ensure the consistency of the shared resources .
- ▶ **Push-through blocking.** It occurs when a medium-priority job is blocked by a lower-priority job that has inherited a higher priority from a job it directly blocks.
 - ▶ Push-through blocking is to **avoid** unbounded priority inversion.



Priority Inheritance Blocking Time

- ▶ Whenever a high priority job wants a resource locked by a lower priority job, the lower priority job gains the priority of the higher priority job as long as the higher priority job is waiting.
- ▶ The high priority job only needs to wait for one additional blocking time.
- ▶ So, we need to know the maximum time any lower priority job will hold a lock on the resource.



Blocking Properties

- ▶ **Lemma 7.1** A semaphore S_k can cause push-through blocking to job J_i , only if S_k is accessed both by a job with priority lower than P_i and by a job that has or can inherit a priority equal to or higher than P_i .
- ▶ **Lemma 7.3** If there are n lower-priority jobs that can block a job J_i , then J_i can be blocked for at most the duration of n critical sections (one for each of the n lower priority jobs), regardless of the number of semaphores used by J_i .
 - ▶ A job can block another if it is inside a critical section. Once out, it can't enter another since it no longer has a high priority.
- ▶ **Lemma 7.4** If there are m distinct semaphores that can block a job J_i , then J_i can be blocked for at most the duration of m critical sections, one for each of the m semaphores.
 - ▶ A nested critical section uses the outer most duration since it is the longest.
- ▶ **Theorem 7.1 (Sha-Rajkumar-Lehoczky)** Under the Priority Inheritance Protocol, a job J can be blocked for at most the duration of $\min(n, m)$ critical sections, where n is the number of lower-priority jobs that could block J and m is the number of distinct semaphores that can be used to block J .



Finding the Blocking Time

- ▶ For each semaphore s_k we define a *ceiling* $C(s_k)$ to be the priority of the highest priority task that may use it.
- ▶ n is the number of lower-priority *jobs* that could block J , m is the number of distinct *semaphores* that can be used to block J .

Lemma 7.5 *In the absence of nested critical sections, a critical section $Z_{j,k}$ of J_j guarded by S_k can block J_i only if $P_j < P_i \leq C(S_k)$.*

$$B_i^l = \sum_{j=i+1}^n \max_k [D_{j,k} : C(S_k) \geq P_i]$$

$$B_i^s = \sum_{k=1}^m \max_{j>i} [D_{j,k} : C(S_k) \geq P_i].$$



Finding the Blocking Time

$$B_i^l = \sum_{j=i+1}^n \max_k [D_{j,k} : C(S_k) \geq P_i]$$
$$B_i^s = \sum_{k=1}^m \max_{j>i} [D_{j,k} : C(S_k) \geq P_i].$$

1. For each job J_j with priority lower than P_i , identify the set $\beta_{i,j}$ of all critical sections that can block J_i .
 2. For each semaphore S_k , identify the set $\gamma_{i,k}$ of all critical sections guarded by S_k that can block J_i .
 3. Sum the duration of the longest critical sections in each $\beta_{i,j}$, for any job J_j with priority lower than P_i ; let B_i^l be this sum.
 4. Sum the duration of the longest critical sections in each $\gamma_{i,k}$, for any semaphore S_k ; let B_i^s be this sum.
 5. Compute B_i as the minimum between B_i^l and B_i^s .
-



Calculating Blocking Time

Theorem 7.2 *A set of n periodic tasks using the Priority Inheritance Protocol can be scheduled by the Rate-Monotonic algorithm if*

$$\forall i, 1 \leq i \leq n, \quad \sum_{k=1}^i \frac{C_k}{T_k} + \frac{B_i}{T_i} \leq i(2^{1/i} - 1). \quad (7.2)$$

	$S_1(P_1)$	$S_2(P_1)$	$S_3(P_2)$
J_1	1	2	0
J_2	0	9	3
J_3	8	7	0
J_4	6	5	4

$$B_i^l = \sum_{j=i+1}^n \max_k [D_{j,k} : C(S_k) \geq P_i]$$

$$B_i^s = \sum_{k=1}^m \max_{j>i} [D_{j,k} : C(S_k) \geq P_i].$$

$$B_1^l = 9 + 8 + 6 = 23$$

$$B_1^s = 8 + 9 = 17 \quad \Rightarrow B_1 = 17$$

$$B_2^l = 8 + 6 = 14$$

$$B_2^s = 8 + 7 + 4 = 19 \quad \Rightarrow B_2 = 14$$

$$B_3^l = 6$$

$$B_3^s = 6 + 5 + 4 = 15 \quad \Rightarrow B_3 = 6$$

$$B_4^l = B_4^s = 0$$

$$\Rightarrow B_4 = 0$$

Nested Critical Sections




Figure 7.7 Priority inheritance with nested critical sections.

PIP Problems

- ▶ If we allow multiple blockings, we need the B_i for each resource.
- ▶ PIP does not prevent deadlock




Figure 7.11 Example of deadlock.

Priority Ceiling Protocol

- ▶ Fixed priority scheduling
- ▶ Uses priority inheritance
- ▶ For each resource S_i define its **resource priority ceiling** $\pi(S_i) = \max$ priority of all jobs which may lock S_i
- ▶ **System Priority Ceiling** at t , $\hat{\pi}(t) = \max \pi(S_i)$ of all locked resources



PCP Rules

- ▶ Suppose J wants to lock a resource R at time t
 - ▶ If R is held by some other job J_L
 - ▶ J is blocked and J_L should inherit the priority of J
- ▶ If the resource is free,
 - ▶ and if J 's priority is higher than **system ceiling** $\hat{\pi}(t)$, then J can lock the resource
 - ▶ If J 's priority is not higher than $\hat{\pi}(t)$, but if J is holding the resource R_m with $\hat{\pi}(t) = \pi(R_m)$, then J still can lock it
 - ▶ Else, block



PCP Example

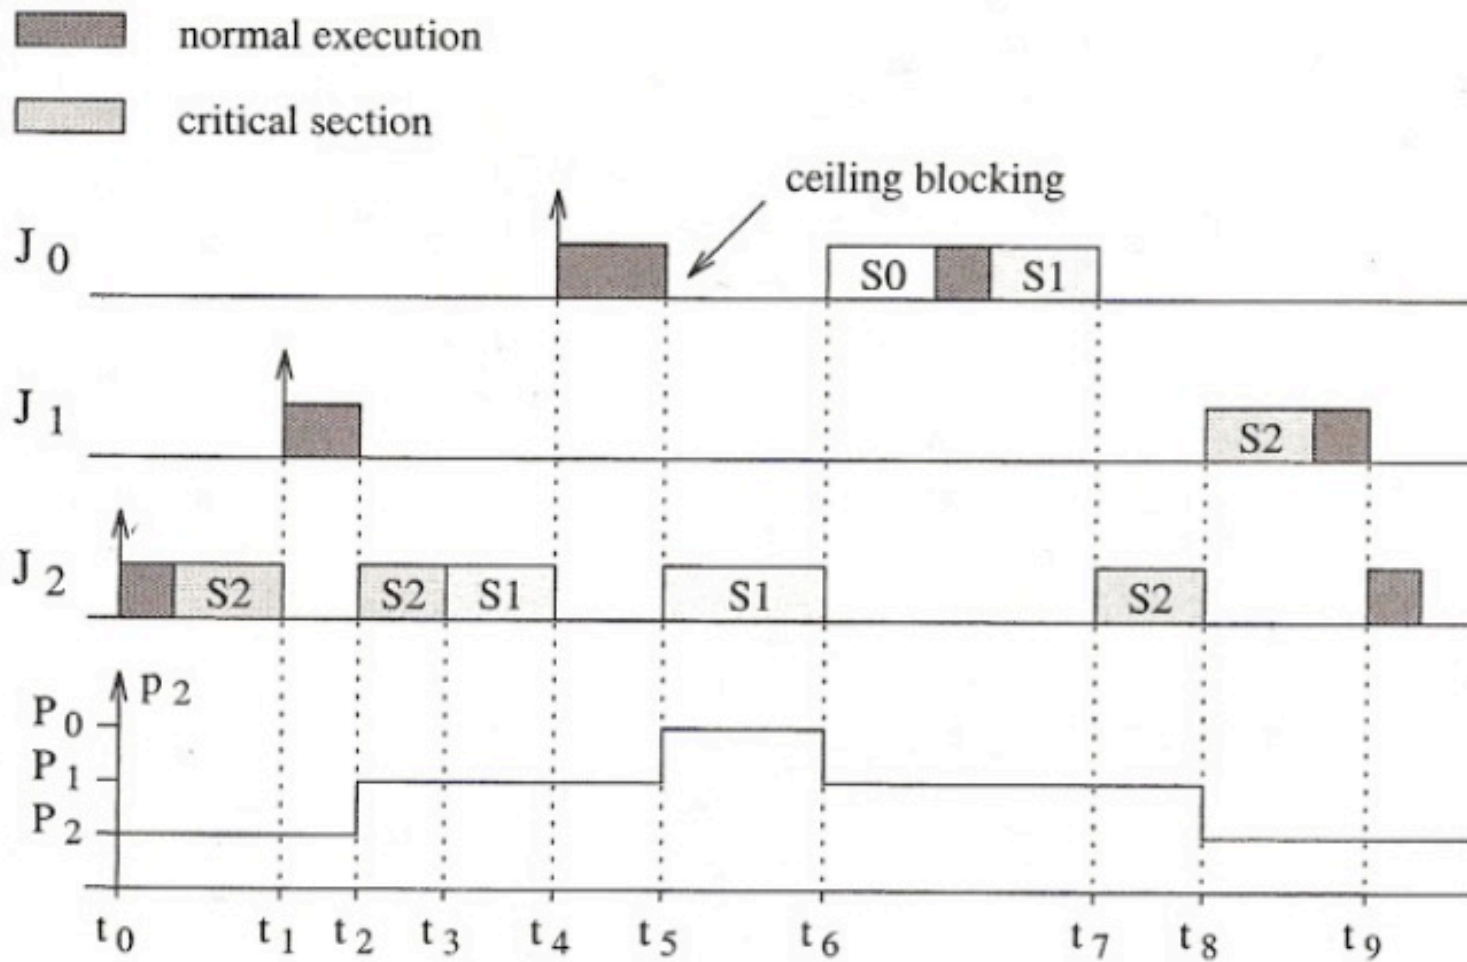



Figure 7.12 Example of Priority Ceiling Protocol.

PCP

- ▶ PCP avoids transitive blocking.
 - ▶ In the earlier deadlock example, J_H will be blocked since the system priority ceiling will be H when J_L locks R_1



- ▶ Wait for any resource that has a higher priority ceiling than you

PCP Properties

- ▶ **Lemma 7.6** *If a job J_k is preempted within a critical section Z_a by a job J_i that enters a critical section Z_b , then, under the Priority Ceiling Protocol, J_k cannot inherit a priority higher than or equal to that of job J_i until J_i completes.*
- ▶ The maximum blocking time B_i of job J_i can be computed as the duration of the longest critical section among those belonging to tasks with priority lower than P_i and guarded by a semaphore with ceiling higher than or equal to P_i .
- ▶ If $D_{j,k}$ denotes the duration of the longest critical section of task T_j among those guarded by semaphore S_k ,

$$B_i = \max_{j,k} \{D_{j,k} \mid P_j < P_i, C(S_k) \geq P_i\}. \quad (7.6)$$



PCP Schedulability

► Using the same example as PIP,

Theorem 7.2 *A set of n periodic tasks using the Priority Inheritance Protocol can be scheduled by the Rate-Monotonic algorithm if*

$$\forall i, 1 \leq i \leq n, \quad \sum_{k=1}^i \frac{C_k}{T_k} + \frac{B_i}{T_i} \leq i(2^{1/i} - 1). \quad (7.2)$$

	$S_1(P_1)$	$S_2(P_1)$	$S_3(P_2)$
J_1	1	2	0
J_2	0	9	3
J_3	8	7	0
J_4	6	5	4

$$\begin{cases} B_1 = \max(8, 6, 9, 7, 5) = 9 \\ B_2 = \max(8, 6, 7, 5, 4) = 8 \\ B_3 = \max(6, 5, 4) = 6 \\ B_4 = 0. \end{cases}$$



PIP vs PCP

- ▶ Priority Ceiling Protocol is more restricted than Priority Inheritance Protocol
- ▶ To lock a resource:
 - ▶ In PIP – Check if a resource is locked
 - ▶ In PCP – Check if job priority is higher than the current system resource priority to lock a resource.



PIP vs PCP Blocking Time

- ▶ PIP blocking time is l per resource or job if you do not allow nesting.
 - ▶ Nested: If you want a resource (A) with someone who has it locked, and they lock (B) as well, need to consider everyone who can lock (B)
- ▶ PCP blocking time is l (the longest critical section).

